A Computational Proof of the Highest-Scoring Boggle Board

Dan Vanderkam

2025-06-08

Abstract Finding all the words on a Boggle board is a classic computer programming problem. With a fast Boggle solver, local optimization techniques such as hillclimbing and simulated annealing can be used to find particularly high-scoring boards. The sheer number of possible Boggle boards has historically prevented an exhaustive search for the global optimum board. We apply Branch and Bound and a tailor-made data structure to perform the first such search. We find that the highest-scoring boards found via hillclimbing are, in fact, the global optima.

Introduction

Boggle is a word search game invented in 1972 by Allan D. Turoff and currently sold by Hasbro. Competitors shake up 16 dice to form a 4x4 grid of letters. The goal is to find as many words as possible in three minutes. Letters can be connected up, down, left, right and diagonally, and words need not be found in a straight line. The same die cannot be used twice in a word. Longer words count for more points (3/4 letters=1 point, 5=2, 6=3, 7=5, 8+=11 points).

Finding all the words on a Boggle board has become a classic computer programming problem. It is often assigned in classes, used as an interview question and, more recently, given as a task to LLMs. With a fast Boggle solver, it's natural to search for particularly high-scoring boards. Typically this is done via hillclimbing, simulated annealing or genetic algorithms. Searches of this kind date back to at least 1982. While these searches do produce high-scoring boards, they cannot make definitive statements about whether these are *the* highest-scoring boards.

This paper takes a different approach. By using Branch and Bound, a data structure and algorithms tailor-made for Boggle, and a large amount of compute, we're able to establish for the first time that the best Boggle boards found via hillclimbing are, in fact, the global optima.

P	E	R	S
L	A	T	G
S	I	N	E
T	E	R	S

Table 1: The highest-scoring Boggle board for the ENABLE2K dictionary, with 1,045 words and 3,625 points. The longest word is "replastering."

Terminology and Conventions

The words that can be found on a Boggle board are determined by the letters on the board and by the choice of dictionary. In this paper, we'll use the ENABLE2K word list, which was developed by Alan Beale in 1997 and 2000. This word list contains 173,528 entries.

We adopt the following terminology and conventions:

- B refers to a Boggle board. To refer to a specific Boggle board, we write out the letters of the board in column-major order. (This distinction is only important for non-square dimensions such as 2x3 and 3x4 Boggle, which lack reflectional symmetry.)
- Because one of the Boggle dice contains a "Qu" (two letters), we adopt the convention that `q` indicates a Qu cell. So `qaicdrneetasnnil` refers to the board in Figure N.
- Boggle dice use uppercase letters (except for Qu), but we typically use lowercase. No meaningful distinction

is drawn between uppercase and lowercase in this paper.

- The cells on an MxN board are numbered $0...M * N - 1$ in column-major order, as shown in Figure N. We refer to the letter on cell i of board B as B_i.
- Words(B) is the set of all words that can be found on the board B.
- Score(B) is the sum of the point value of these words, $\sum_{w \in \text{ Words}(B)} \text{SCORES}[\text{Len}(w)]$. We refer to this as the score of the board.

Heuristics to find high-scoring boards

Finding all the words on a Boggle board is a popular programming puzzle. This is done by performing a depth-first search over the board's adjacency graph starting at each cell. The key to making this efficient is to prune out prefixes such as "bnj" that don't begin any words in the dictionary. This is typically done using a Trie (Prefix Tree) data structure.

```
# Listing 0: Scoring a Boggle Board
def score(bd: str, trie: Trie) -> int:
  used = {}

  def step(idx: int, node: Trie) -> int:
    score = 0
    used[idx] = True
    if node.has_child(bd[idx]):
      n = node.child(bd[idx])
      if n.is_word() and not n.is_visited():
        score += SCORES[n.length()]
        n.set_visited()
      for n_idx in NEIGHBORS[idx]:
        if not used.get(n_idx):
          score += step(n_idx, n)
    used[idx] = False
    return score

  score = 0
  for i in range(m * n):
    score += step(i, trie)
  return score
```

Nodes in the Trie are marked as we find words to avoid double-counting. With some care, it is possible to find the score of individual boards extremely rapidly on modern hardware. For example, the author's laptop is able to score random 4x4 Boggle boards at a rate of around 200,000 boards per second.

This speed can be used to attack a new problem: finding high-scoring boards. This is typically done via local search heuristics such as hillclimbing, simulated annealing, or genetic algorithms. A particularly effective approach for Boggle is to iteratively explore around a pool of high-scoring boards, as shown in Listing 1.

```
# Listing 1: Hillclimbing Algorithm
N = 500  # pool size
Pool <- N random Boggle Boards
Repeat until convergence:
  Neighborhood <- Pool + Boards within edit
distance 1 of a board in Pool
  Pool <- N highest-scoring boards in
Neighborhood
```

We take an "edit" to mean changing one letter or swapping two letters. With a pool size of N=500, this usually (95/100 times) converges with the highest-scoring board as `perslatgsineters` (Figure N), which contains 1,045 words and scores 3,625 points using ENABLE2K.

This makes one wonder whether this board is, in fact, the global optimum. The principal result of this paper is to show that it is.

Proof of Maximality

The most straightforward way to prove that a board is the highest-scoring is by exhaustive search. Unfortunately, the combinatorial explosion of possible boards renders this infeasible for all but the smallest sizes:

Dims	Boards	Rate	Time
3x3	26^9/8 ≈ 6.8e11	600k bd/s	~12 days
3x4	26^12/4 ≈ 2.4e16	400k bd/s	~2000 years
4x4	26^16/8 ≈ 5.5e21	200k bd/s	~900M years

One objection is that not all 26^16 combinations of letters can be rolled with the standard 16 Boggle dice. Determining whether a particular letter combination can be rolled is a Set-Cover problem, which is NP-Complete. A greedy approach works well for this small problem, however, and we can estimate that approximately 1 in 12 combinations of letters can be rolled. While this reduces the search space, it's not enough to make exhaustive search feasible, and this will not prove to be a useful constraint.

Two observations suggest a path towards a solution:

1. Similar boards have similar sets of words on them. Finding the score of a board and that of its neighbors involves large amounts of duplicated work.

2. Most boards have scores that are significantly lower than the best one. The average score of a random board is about 42, nearly 100x less than 3625. (The average score of a board rolled with Boggle dice is closer to 120.)

The first observation suggests that we should group boards together to reduce repeated work. The second indicates that we have considerable "wiggle room" to upper bound the score rather than calculate it precisely.

Branch and Bound

Rather than exhaustive search, we use Branch and Bound to find the globally optimal board. Branch and Bound is an algorithm design paradigm dating back to the 1960s that narrows in on the global optimum by recursively subdividing the search space.

To apply it, we need to define two operations on sets of Boggle boards:

- bound(S): an upper bound on the score of any board in a set of boards S
- branch(S): a way to split the set S into smaller sets S_1, S_2, ..., S_m.

With these operations in place, the Branch and Bound algorithm to find all boards with a score greater than S_{high} is given in Listing 2:

```
# Listing 2 - Branch and Bound Algorithm
Queue <- {Universal Set of MxN Boggle Boards}
while Queue is non-empty:
  S <- Pop(Queue)
  if |S| == 1:
  S is a candidate solution; score its lone
board to check
  else if bound(S) < S_high:
  S cannot contain a high-scoring board; it
can be dropped.
  else:
  for S_i in branch(S):
    Queue <- S_i
```

The appeal of this approach is that, when bound(S) is low, we can discard the entire set S without having to evaluate every board in it. Now we need to define the branch and bound operations.

Board classes and the branch operation

Rather than allowing arbitrary sets of Boggle boards, we restrict ourselves to "board classes." These require each cell of a board in the set to come from a particular set of letters:

$$C(L_1, L_2, ..., L_{m \cdot n}) = \{B \mid B_i \in L_i \forall i\}$$

For example, here's a 3x3 board class where each cell can be one of two letters:

{A,B}	{G,H}	{M,N}
{C,D}	{I,J}	{O,P}
{E,F}	{K,L}	{Qu,R}

Table 3: A 3x3 board class containing 512 boards

This board class contains 2^9 = 512 possible boards. Here are a few of them:

A	G	M
D	I	P
E	L	R

agmdipelr: 85 points

B	H	M
C	I	P
E	L	R

bhmcipelr: 62 points

B	H	M
D	J	P
F	K	Qu

bhmdjpfkq: 0 points

Analogous to the B_i notation for boards, we can indicate the possible letters on a cell in a board class as C_i. On this board class, for example, C_0 = `{"A", "B"}`.

We can carve the 26 letters of the alphabet up into distinct "buckets" to reduce the combinatorial explosion of possible boards into a much smaller number of board classes.

The following partitions of the alphabet were found via a heuristic search:

N	Letter Buckets
2	{aeiosuy, bcdfghjklmnpqrtvwxz}
3	{aeijou, bcdfgmpqvwxz, hklnrsty}
4	{aeiou, bdfgjqvwxz, lnrsy, chkmpt}

Using three buckets with 4x4 Boggle, for example, we have only $\frac{3^{16}}{8} \approx 5.4 * 10^6$ board classes to consider, a dramatic reduction from the $5.5 * 10^{21}$ individual

boards. Of course, this will be in vain if operations on board classes are proportionally slower. Fortunately, this will not prove to be the case.

To "branch" from a board class, we split one of its cells into the individual possibilities. For example, starting with this board class containing 5,062,500 boards:

lnrsy	chkmpt	lnrsy
aeiou	aeiou	aeiou
chkmpt	lnrsy	bdfgjqvwxz

we might split the center cell to get five board classes containing 1,012,500 boards each:

lnrsy	chkmpt	lnrsy
aeiou	a	aeiou
chkmpt	lnrsy	bdfgjqvwxz

lnrsy	chkmpt	lnrsy
aeiou	e	aeiou
chkmpt	lnrsy	bdfgjqvwxz

...

lnrsy	chkmpt	lnrsy
aeiou	u	aeiou
chkmpt	lnrsy	bdfgjqvwxz

Since the center and edge cells have the greatest connectivity, we split these first before splitting the corners.

Board classes still have all the same symmetries as a Boggle board. This allows us to only consider "canonically-oriented" board classes for a roughly 8x reduction in the search space.

The sum bound

Next we need to construct an upper bound. One possible bound is the score of every word that can appear on any board in the board class.

$$\text{sum_bound}(C) = \sum_{B \in C} \sum_{\text{Words}(B)} \text{SCORES[w]}$$

Here the sums are taken over *unique* words. This can be calculated in a similar manner to an ordinary Boggle solver, except that we need two loops now: one for neighbors, and a new one for each possible letter on each cell:

```
# Listing 3: sum bound on a Boggle board class
def sum_bound(
  board_class: list[str], trie: Trie
) -> int:
  used = {}

  def step(idx: int, node: Trie) -> int:
    score = 0
    used[idx] = True
    letters = board_class[idx]
    for letter in letters:  # new loop
      if node.has_child(letter):
        n = node.child(letter)
        if n.is_word() and not n.is_visited():
          score += SCORES[n.length()]
          n.set_visited()
        for n_idx in NEIGHBORS[idx]:
          if not used.get(n_idx):
            score += step(n_idx, n)
    used[idx] = False
    return score

  score = 0
  for i in range(m * n):
    score += step(i, trie)
  return score
```

Clearly we have

$$\text{sum_bound}(C) \geq \text{Score}(B) \forall B \in C$$

because every word on every possible board contributes to the bound.

A useful property of this bound is that, if `C = {B}`, then `sum_bound(C) = Score(B)`, that is to say, it converges on the true Boggle score for single-board sets.

Unfortunately, this bound is imprecise because it doesn't take into account that some letter choices are mutually exclusive. For example, consider this 2x2 board class containing two individual boards:

F	{A, U}
.	R

There are two words here, "far" and "fur." These each count for 1 point, so the sum bound of the board class is 2. No individual board can contain both of these words, however, since the A and the U are mutually exclusive, so this is an overestimate.

This proves problematic for large board classes. For example, the sum_bound of the 5,062,500 board class

above is 109,524 points, but the best board in that class only scores 545 points. To get a better bound, we need to take into account that choices on cells are exclusive.

The max bound

We can model this by taking the the max across the letter possibilities on a cell instead of the sum. In doing so, we dispense with any attempt to enforce the constraint that a word can only be found once.

```
# Listing 4: max bound on a Boggle board class
def max_bound(
  board_class: list[str], trie: Trie
) -> int:
  used = {}

  def step(idx: int, node: Trie) -> int:
    score = 0
    used[idx] = True
    letters = board_class[idx]
    for letter in letters:
      if node.has_child(letter):
        letter_score = 0
        n = node.child(letter)
        if n.is_word():
          letter_score += SCORES[n.length()]
        for n_idx in NEIGHBORS[idx]:
          if not used.get(n_idx):
            letter_score += step(n_idx, n)
        score = max(score, letter_score)
    used[idx] = False
    return score

  bound = 0
  for i in range(m * n):
    bound += step(i, trie)
  return bound
```

We can see that this is a valid bound because, for any particular board `B` in a class `C`:

1. It produces the full set of recursive calls for `B` from Listing 0, as well as many other calls.
2. For each of these matching calls, `step` returns a score greater than or equal to the `step` call in `score`. This could be either because there's another letter choice that produces a higher score, or because `max_bound` double-counts a word that `score` does not.

So we have

$$\text{max_bound}(C) \geq \text{Score}(B) \forall B \in C$$

In practice, this bound is considerably tighter than the sum bound (see Table N). However, because it double-counts words, the max bound for a board class containing a single board may be greater than the score of that board. (This can only happen if the board contains a repeated letter.)

Here are the sum and max bounds for the 5,062,500 board 3x3 class and each of its five splits:

Center Cell	Sum Bound	Max Bound	True Max
aeiou	109,524	9,460	545
a	56,576	6,120	545
e	72,026	7,023	520
i	60,244	6,231	503
o	49,533	5,525	392
u	38,214	4,464	326

The max bound is an order of magnitude tighter than the sum bound in all cases. It's still imprecise, however, because it might choose different letters for a cell along different search paths in the DFS. Consider this 2x2 board class:

T	I
{A, E}	R

Starting with the "T:"

- If we go down, we can form "TAR" by picking the "A" but we cannot form any words if we pick the "E." So the "T" → {A, E} path in the `max_bound` DFS nets 1 point.
- If we go right to "I", we can only score points by picking the "E" for this cell to form "TIE" and "TIER." So the "T" → "I" → {A, E} path nets 2 points.

The points from these two paths are added. But no single board can have both an "A" and an "E" on the bottom left cell, so neither board in this class contains both "TAR" and "TIE." This is why the max bound is imprecise. It enforces that we make a choice on each cell, but not that this choice be consistent across all paths through that cell.

The minimum of two upper bounds is also an upper bound, so we can also use:

$$\text{max_sum_bound}(C) = \min(\text{max_bound}(C), \text{sum_bound}(C))$$

as an upper bound that combines the strengths of both. These bounds can be calculated simultaneously in a single DFS.

Initial Results with Branch and Bound

Using the Branch and Bound algorithm with board classes and `max_sum_bound` results in a dramatic speedup over exhaustive search. For 3x3 Boggle using three buckets on the author's laptop, the search completes in about an hour on a single CPU core. This represents roughly a 300x speedup. The highest-scoring 3x3 boards found via Branch & Bound precisely match those found via hillclimbing.

S	T	R
E	A	E
D	L	P

Table 12: The highest-scoring 3x3 board, with 545 points. Long words include "repasted" and "replated."

This speedup makes 3x3 Boggle maximization easy on a laptop and 3x4 maximization possible in a data center. But it offers little hope for 4x4 Boggle.

Despite the speedup, there remains an enormous amount of repeated work. Each evaluation of max_sum_bound is performed independently, but the computation for max_sum_bound(C) and its children after the "branch" operation (max_sum_bound(C1), max_sum_bound(C2), ...) is nearly identical. To achieve a greater speedup, we'll seek to eliminate this repetition.

Sum/Choice trees

Our goal is to speed up repeated branch and bound calculations. To do so, we'll forget about `sum_bound`, whose global uniqueness is difficult to maintain. Instead, we'll focus solely on `max_bound`, which can be be more easily calculated using local information.

Previously `max_bound` was calculated using recursive function calls. Our next step is to convert these function calls into a tree structure in memory. This will allow us to implement branch and bound as operations on the tree.

First, we refactor `max_bound` to use two functions. These will become two types of nodes in our tree:

```
# Listing 5: max bound with two functions
def max_bound(
  board_class: list[str], trie: Trie
) -> int:
  used = {}

  def choice_step(idx, node):
    score = 0
    used[idx] = True
    letters = board_class[idx]
    for letter in letters:
      if node.has_child(letter):
        child = node.child(letter)
        score = max(
          score,
          sum_step(idx, child),
        )
    used[idx] = False
    return score

  def sum_step(idx, node):
    score = 0
    if node.is_word():
      points = SCORES[node.length()]
      score += points
    for n_idx in NEIGHBORS[idx]:
      if not used.get(n_idx):
        score += choice_step(n_idx, node)
    return score

  bound = 0
  for i in range(m * n):
    bound += choice_step(i, trie)
  return bound
```

This is a simple transformation of the previous `max_bound`. With this new formulation, we construct a tree where each node corresponds to one of the function calls:

```
Node := SumNode | ChoiceNode

ChoiceNode:
  cell: int
  children: {letter -> SumNode}

SumNode:
  points: int
  children: ChoiceNode[]
```

Note that `points` is the points on an individual node, not the bound for the entire subtree. The top-level call to `max_bound` can be modeled as a `SumNode` with each cell as a child:

$$\text{BuildTree}(C) \rightarrow \text{SumNode}$$

A direct translation of the call graph results in numerous "dead paths" that do not lead to any points. These can be pruned to produce a more compact tree.

The bound for each node can be computed as:

```
Bound(n: SumNode)
= n.points + sum(Bound(c) for c in n.children)
```

```
Bound(n: ChoiceNode)
= max(Bound(c) for c in n.children)
```

In practice, the bound can be stored explicitly on each node and updated as we modify the tree. Here's what one of these trees looks like for the TAR/TIER board from earlier:

cell 0: T	cell 2: I
cell 1: {A, E}	cell 3: R

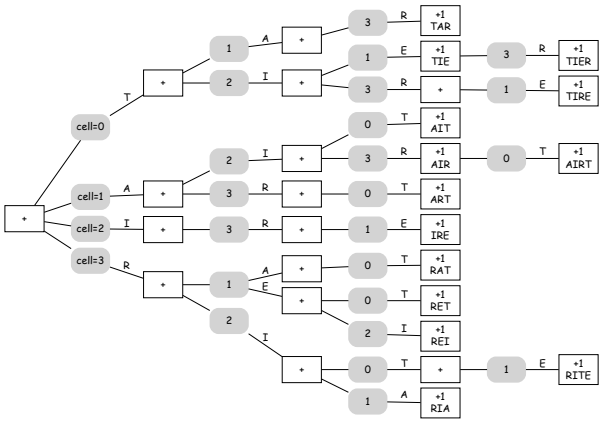

Figure 1: Tree for the TAR/TIER board class. SumNodes are rectangular, ChoiceNodes are round. Choices of letters on ChoiceNodes are indicated along edges. Some of these words (AIT, AIRT, REI, RET) are obscure, but are valid Boggle plays.

Here's the same tree showing the bound on each node:

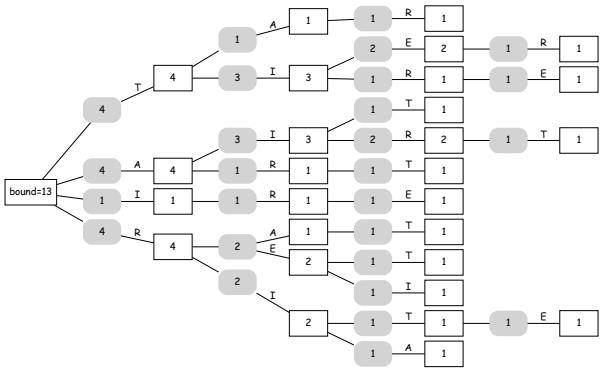

Figure 2: Same tree showing `Bound` on each node.

We can make a few observations about these Sum/Choice trees:

- By construction, `Bound(BuildTree(C))` = `max_bound(C)`.
- The wordlist and geometry of the Boggle board are fully encoded in the tree. Once the tree is constructed, we no longer need to reference the Trie or the `NEIGHBORS` array.
- Words correspond to `SumNodes` with points on them. A `SumNode` has zero or one words associated with it.
- Individual words can be read off by descending the tree and tracking the letters used on each `ChoiceNode`.
- `ChoiceNodes` for the same cell may appear multiple times in the tree. The bound is imprecise because the `max` operation may not make the same choice on each `ChoiceNode`.

Multiboggle and the Invariant

We've seen that the root bound on the tree is an alternate way to calculate `max_bound` for a board class. Now we want to perform operations on these trees, and these operations may affect the bound. To prove that the bound remains valid, we'll establish an invariant that implies the validity of the bound. Then we'll show that each operation maintains this invariant.

First, we define the "Force" operation on a tree:

```
Force(n: SumNode, B) = n.points + sum(Force(c,
B) for c in n.children)
Force(n: ChoiceNode, B) =
Force(n.choices[B_{n.cell}], B) or 0
```

Intuitively, this "forces" each cell to match the board `B`.

Lemma: If $T = \text{BuildTree}(C)$, then

$$\text{Force}(T, B) \leq \text{Bound}(T) \forall B \in C$$

This is immediate from the definition. Force is the same as Bound on Sum nodes, and less than or equal to Bound on Choice nodes.

So what is $\text{Force}(T, B)$? We can write out code to calculate this by modifying Listing N:

```python
# Listing 6: Build+Force operation on Tree
def forced_tree(
  board_class: list[str], board: str, trie
):
  used = {}

  def choice_step(idx, trie):
    score = 0
    used[idx] = True
    letter = board[idx]
    if trie.has_child(letter):
      n = trie.child(letter)
      score = sum_step(idx, n)
    used[idx] = False
    return score

  def sum_step(idx, trie):
```

```
    score = 0
    if trie.is_word():
      score += SCORES[trie.length()]
    for n_idx in NEIGHBORS[idx]:
      if not used.get(n_idx):
        score += choice_step(n_idx, trie)
    return score

  bound = 0
  for i in range(m * n):
    bound += choice_step(i, trie)
  return bound
```

We can make a few immediate observations:

1. `Force(T, B)` does not depend on the board class `C`. It is a function of `B` alone.
2. `Force(T, B)` performs the exact same calculation as `Score(B)`, except that there are no checks for whether a word has been found more than once.

We'll refer to this as Multi(B), the "Multiboggle score" of B. This can be thought of as a variation on Boggle where you're allowed to find the same word multiple times. For example, this 2x3 Boggle board:

```
E B E
E F E
```

Has Score(B) = 3 ("bee", "fee", "beef") but Multi(B) = 12 because each word can be found along four distinct paths.

- **Lemma**: Multi(B) $\geq$ Score(B). This is clear from the definition. The Multiboggle score is an upper bound on the Boggle score.
- **Lemma**: Multi(B) = Score(B) if B does not contain repeated letters. (The converse is not true.)
- **Lemma**: max_bound($\{B\}$) = Multi(B).

In other words, the `max_bound` converges to the Multiboggle score as you progressively force cells on a board class.

Putting this together, if T = BuildTree(C) then we have:

$$\begin{aligned} \text{Score}(B) &\leq \text{Multi}(B) \\ &= \text{Force}(T, B) \\ &\leq \text{Bound}(T) \quad \forall B \in C \end{aligned}$$

So if we can show that Force(T, B) = Multi(B) for all boards in a board class, then Bound(T) is a valid upper bound for Score(C).

For most boards, Multi(B) is close to Score(B). Since we have considerable "wiggle room" between the average score of a board (~40 points) and the score of the best board (3625), working with the Multiboggle score is usually an acceptable concession. What we'll seek is boards B with Multi(B) $\geq S_{\text{high}}$. For each of these, we can confirm whether Score(B) $\geq S_{\text{high}}$ as well using a regular Boggle solver.

While Multi(B) is usually close to Score(B), there are some pathological cases where this breaks down. For example, the board in Figure N has Score(B) = 189, but Multi(B) = 21953! (The word "reservers" can be found in 100 distinct ways.) We will partially address this issue later in the paper.

E	E	E	S
R	V	R	R
E	E	E	S
R	S	R	S

Table 14: Pathological board with Score(B)=189 but Multi(B)=21,953.

Sum/Choice Satisfiability is NP-Hard

We seek boards B in a board class C such that Force(T, B) $\geq S_{\text{high}}$. Since each board in a board class represents a choice of letter on each of the cells, we can think of this as a satisfiability problem.

Theorem: Determining whether there exists B such that Force(T, B) $\geq S_{\text{high}}$ is NP-Hard.

Proof: We map from 3-CNF, a known NP-Hard problem, to the Sum/Choice Tree satisfiability problem.

Suppose we have a 3-CNF formula with m clauses on $x_1, x_2, ...x_n$.

For each clause, we construct a tree which evaluates to 1 if the clause is satisfied and zero if it is not satisfied.

- If the clause contains a single term a, then we model this as a ChoiceNode on cell a.
- If the clause is $a \vee b$, we model is as a tree with two layers of `ChoiceNodes`.
- If the clause is $a \vee b \vee c$, we model it as a tree with three layers of `ChoiceNodes`.

Finally, we create a root `SumNode` T with the m `ChoiceNodes` as children. By construction, $\exists B \mid \text{Force}(T, B) = m$ iff there are x_i that satisfy the 3-SAT problem. So if we can solve the satisfiability problem for Sum/Choice trees, we can also solve it for 3-

CNF. Since 3-CNF is known to be NP-Hard, this means that Sum/Choice satisfiability is NP-Hard as well.

(source: D.W.'s mapping on Stack Overflow)

So we should not expect to find an efficient solution to this problem, nor one that scales well to larger boards. This doesn't necessarily mean that Boggle maximization itself is NP-Hard, since not every Sum/Choice tree corresponds to a Boggle board. Still, it is suggestive that this is a hard problem.

Orderly Trees

Before defining operations on general Sum/Choice trees, it will be helpful to shift our perspective on them. So far, we've thought of them as tree representations of the recursive call structure of `max_bound`.

An alternative view, however, is to treat them as a container structure holding paths to words and the points associated with those words. Instead of forming the tree via DFS, we can find all the paths to words in the board class and add each of them to the tree structure.

We can define a path as a sequence of cells and letters on those cells:

```
Path p = list((cell, letter))
```

Then we can define `add_word`:

```
# Listing 7: add_word to sum/choice tree
def add_word(
  node: SumNode, path: Path, points: int
):
  if len(path) == 0:
    node.points += points
    return node

  cell, letter = path[0]
  choice_n = find(
    node.children, lambda c: c.cell == cell
  )
  if not choice_n:
    choice_n = ChoiceNode(cell=cell)
    node.children.append(choice_n)

  sum_n = choice_n.children.get(letter)
  if not sum_n:
    sum_n = SumNode()
    choice_n.children[letter] = sum_n

  return add_word(sum_n, path[1:], points)
```

Lemma: This produces an identical tree.

Every path to points is present and identical in both constructions.

This shift in perspectives allows us to establish the critical result:

Theorem: Anagramming words before adding them preserves the invariant.

This can be seen by treating the `Force` operation as a sum across all SumNodes with points in the tree, conditioned on whether the path to that node is realized in the board. We can define the path to a node by tracing it back to the root:

```
Path(n: SumNode) =
  [] if n is the root
  Path(n.parent.parent) ++ [(n.parent.cell,
n.letter)]
```

Then we can reformulate `Force` as a sum across compatible paths:

$$\text{Force}(T, B) =$$
$$\sum\big(\text{node.points } \forall \text{ node} \in T$$
$$\mid \text{AND}_{n \in \text{ Path(node)}}(B_{n.\text{cell}} = n.\text{letter})\big)$$

The anagramming theorem is a natural consequence of `AND` being commutative.

Previously, words were added to the tree in the order in which they were spelled. We can see now, however, that this was a choice. To get more consistent ordering, and thus lower bounds, we can define a canonical order for the cells and sort the paths to words accordingly before adding them to the tree.

Any ordering is valid but, since center cells are likely to be used in the most words, it makes the most sense to put them at the top of the tree. For 4x4 Boggle, we use the following ordering (higher numbered cells appear closer to the root of the tree):

3	7	5	2
11	15	13	10
9	14	12	8
1	6	4	0

We can produce a tree using this ordering:

```
# Listing 8: Building an orderly tree
def build_orderly_tree(board_class, trie):
  root = SumNode()

  def choice_step(idx, trie, choices):
    letters = board_class[idx]
    for letter in letters:
      if trie.has_child(letter):
        choices.append((idx, letter))
        child = trie.child(letter)
        sum_step(idx, child, choices)
        choices.pop()

  def sum_step(idx, trie, choices):
    if trie.is_word():
      ordered_choices = sorted(
        choices, key=lambda c: -ORDER[c[0]]
      )
      score = SCORES[trie.length()]
      add_word(root, ordered_choices, score)
    for n_idx in NEIGHBORS[idx]:
      if n_idx not in (c[0] for c in choices):
        choice_step(n_idx, trie, choices)

  for i in range(m * n):
    choice_step(i, trie, [])
  return root
```

Because the cells follow a particular order and the resulting tree looks more "well-ordered," we refer to these as Orderly Sum/Choice Trees or just "Orderly Trees."

Using a canonical order for the cells naturally synchronizes choices across subtrees, particularly the choices with a high index (the center). This typically results in smaller trees with lower bounds, especially for large board classes.

Here's the orderly tree for the TAR/TIER board class from earlier, ordered by cell number:

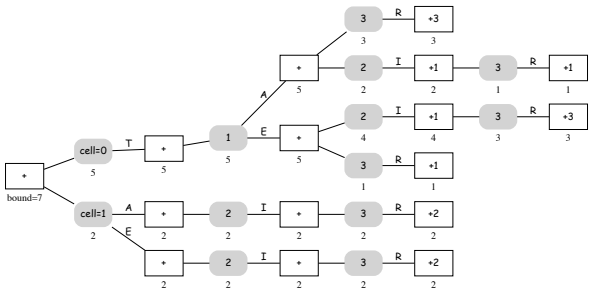

Figure 3: Orderly Tree for "t ae i r"; `Bound(node)` is marked under each node.

Note that the tree is smaller (56⟶29 nodes) and the bound on the root node has dropped from 13 to 7. Before, there were 8 ChoiceNodes with cell=1, but now there are only 2.

Using an Orderly Tree helped for this small board class, but the effect is more dramatic for larger board classes:

Board	max_bound	Orderly bound
2x2	13	7
3x3 (a)	6,361	503
3x3 (b)	9,460	1,523
3x4 (a)	51,317	4,397
3x4 (b)	194,425	10,018
3x4 (c)	69,889	4,452
4x4 (a)	176,937	11,576
4x4 (b)	514,182	53,037

By construction, no ChoiceNode n in an Orderly Tree will have `ORDER[n.cell]` greater than any of its parents. So if a path begins 14⟶6 in the tree, then it may only continue to cells 1 or 4, since those are lower numbers. (14⟶6⟶9 is a valid path, but it would be added to the tree as 14⟶9⟶6.) Intuitively, if cell A has a child on cell B, then this represents all the paths through the board that skip cells between A and B in the canonical order.

We can use this intuition to define Orderly(N):

```
Orderly(n: Int)
Orderly(0) = SumNode with no children
Orderly(N) = SumNode with children:
         OrderlyChoice(i)(cell=i)
         for i = 0 .. (N-1)

OrderlyChoice(n: Int)
Orderly(N) = ChoiceNode with Orderly(N)
children
```

We can make a few more observations about Orderly Trees:

- The Orderly Tree for a board class, and hence the "orderly bound," is dependent on the order of cells that we choose.
- The tree no longer bears any resemblance to a plausible DFS of the board.
- We can no longer associate SumNodes with single words. For example, the "+3" node on the top right of the tree visualization includes the words TAR, RAT and ART. If you can find one of these, you can find all of them. (This is one reason that the Orderly Tree

uses fewer nodes.) Enforcing that each word is only found once would be impossible in this context, since we're not even sure which words we've found.

OrderlyMerge

With Orderly Trees defined, we're finally ready to perform operations on them. Our first goal will be to speed up the "branch" operation and calculation of the subsequent bound. This requires forcing a single cell in the board class to be each of its possible letters and constructing the resulting Orderly Tree.

In practice, it makes the most sense to force the top choice in the tree, i.e. the one with the first position in the canonical order. This is a "merge" operation on Orderly Trees, which is straightforward to implement as in Listing N.

```
# Listing 9: merge operation on orderly trees
# merge: (Orderly(N), Orderly(N)) ->
Orderly(N)
def merge(a: SumNode, b: SumNode) -> SumNode:
  by_cell = {c.cell: c for c in a.children}
  for bc in b.children:
    ac = by_cell.get(bc.cell)
    by_cell[bc.cell] = (
      merge_choice(ac, bc) if ac else bc
    )
  ch = sorted(
    by_cell.values(),
    key=lambda c: -ORDER[c.cell],
  )
  return SumNode(
    points=a.points + b.points, children=ch
  )

# merge_choice: (OrderlyChoice(N),
OrderlyChoice(N))
#               -> OrderlyChoice(N)
def merge_choice(
  a: ChoiceNode, b: ChoiceNode
) -> ChoiceNode:
  ch = {**a.children}
  for choice, bc in b.children.items():
    ac = ch.get(choice)
    ch[choice] = merge(ac, bc) if ac else bc
  return ChoiceNode(cell=a.cell, children=ch)
```

Lemma: $\text{Force}(\text{merge}(T_1, T_2), B) = \text{Force}(T_1, B) + \text{Force}(T_2, B) \forall B \in T_1, T_2$

With the "merge" helper, we can define the "branch" operation. Note that `top_cell` here isn't really a free parameter; if the tree is $\text{Orderly}(N)$, then we must have $\text{top_cell} = N$.

```
# Listing 10: branch operation on orderly
trees
# branch: Orderly(N) -> list(Orderly(N-1))
def branch(
  o: SumNode,
  N: int,
  board_class: list[str],
) -> list[SumNode]:
  top_choice = find(
    o.children, lambda c: c.cell == N
  )
  if not top_choice:
    # cell is irrelevant; o is Orderly(N-1)
    return [o for _ in board_class[N]]

  other_choices = [
    c for c in o.children if c.cell != N
  ]
  skip_tree = SumNode(
    children=other_choices, points=o.points
  )  # Orderly(N-1)
  return [
    merge(
      top_choice.children[letter], skip_tree
    )  # both are Orderly(N-1)
    if top_choice.children.get(letter)
    else skip_tree  # dead letter on cell N.
    for letter in board_class[N]
  ]
```

The `branch` function splits the tree into two parts: one that includes the words that go through the "top" cell (`top_choice`) and another (`skip_tree`) that includes all the words that don't. Each child of `top_choice` corresponds to a particular choice of letter on that cell. Every word must fall into one of these two groups (goes through the cell or doesn't). Adding their bounds will produce a valid bound for this choice of letters. Since the `merge` operation preserves the invariant, the resulting tree will have a valid bound.

TODO: a visual would convey the intuition here, that "branch" is just a merge.

Calling `branch` is considerably faster than building a new tree for each letter choice on a cell. For example, on the high-scoring 4x4 board class from Table N, `branch` split a center cell containing 12 letters and returned 12 subtrees in 0.07s. Building the same trees from scratch took 4.0s, roughly a 60x difference.

A few observations:

- The `merge` operations perform a deep merge.

- In practice we can update a `bound` property on all nodes as we merge them, so that calculating the bound is instant.
- The `merge` operation (and `branch`) are likely to reduce the bound because they synchronize choices across previously distinct subtrees.
- The subtrees we get back from `branch` aren't exactly the same as what you'd get by building orderly trees for the smaller board classes directly. This is because `branch` effectively removes a cell from the board, rather than replacing it with a one-letter choice.

Here are the results of the first `branch` call on the Orderly Tree for the 5 million board 3x3 board class from before.

Center Cell	Nodes	Bound	True Max	max_bound
aeiou	333,492	1,523	545	9,460
a	86,420	1,198	545	6,120
e	98,585	1,417	520	7,023
i	81,062	994	503	6,231
o	75,474	862	392	5,525
u	60,457	753	326	4,464

OrderlyBound

The `branch` operation on its own is sufficient to implement Branch and Bound with Orderly Trees. Eventually we'll merge all the way down to a single `SumNode`, which represents a candidate board. This allocates additional nodes, however, and we may wish to save RAM by not doing that all the way to single boards.

Instead, we can define an alternative algorithm, `OrderlyBound`, which refines the bound on a tree by traversing it, rather than merging subtrees. We'll maintain a stack of pointers to `ChoiceNodes`. To "branch," we'll pop off the ChoiceNodes for the next cell and, for each choice, push all the next child cells. We can maintain a bound as we do this. If the bound ever drops below `S_high`, we can abandon this search path.

```
# Listing 11: orderly_bound
# Assumes N >= 1
def orderly_bound(
  root: SumNode,  # Orderly(N)
  board_class: list[str],
  S_high: int,
):
  def step(
    points: int,
    idx: int,
    # letters chosen on previous cells
    choices: list[char],
    stack: list[ChoiceNode],
  ):
    b = points + sum(bound(n) for n in stack)
    if b < S_high:
      return  # This path is eliminated
    if idx == N:
      # complete board that can't be
eliminated
      record_candidate_board(choices, b)
      return

    # Try each letter on the next cell.
    cell = CELL_ORDER[idx]
    for letter in board_class[cell]:
      next_nodes = [
        n for n in stack if n.cell == cell
      ]
      next_stack = [
        n for n in stack if n.cell != cell
      ]
      next_points = points
      next_choices = choices + [letter]
      for node in next_nodes:
        letter_node =
node.children.get(letter)
        if letter_node:
          next_stack += letter_node.children
          next_points += letter_node.points

      step(
        next_points,
        idx + 1,
        next_choices,
        next_stack,
      )

  step(root.points, 0, [], root.children)
```

Lemma: Each `step` call preserves the invariant that

$$\text{points} + \sum_{n \in \text{ stack}} \text{Force}(n, B) = \text{Multi}(B)$$

for all boards B compatible with `choices`.

The proof is by induction.

$$
\begin{aligned}
&\text{next_points} + \sum_{n \in \text{ next_stacks}} \text{Force}(n, B) \\
&= \text{next_points} + \sum_{n \in \text{ stacks}} \sum_{\text{child } \in n.\text{child[letter]}.\text{children}} \text{Force}(\text{child}, B) \\
&= \text{points} + \sum_{n \in \text{ stacks}} n.\text{child[letter]}.\text{points} + \sum_{n \in \text{ stacks}} \sum_{\text{child } \in n.\text{child[letter]}.\text{children}} \text{Force}(\text{child}, B) \\
&= \text{points} + \sum_{n \in \text{ stacks}} n.\text{child[letter]}.\text{points} + \sum_{\text{child } \in n.\text{child[letter]}.\text{children}} \text{Force}(\text{child}, B) \\
&= \text{points} + \sum_{n \in \text{ stacks}} \text{Force}(n.\text{child[letter]}, B) \\
&= \text{points} + \sum_{n \in \text{ stacks}} \text{Force}(n, B) \qquad \text{definition of Force for ChoiceNode} \\
&= \text{Multi}(B)
\end{aligned}
$$

Base case: The initial call to `step` has the root points, plus all the root's children. This matches the definition of Force for a `SumNode`.

Inductive case: From the inductive hypothesis, for each board B, the call to `step` has

$$\text{points} + \sum_{n \in \text{ stack}} \text{Force}(n, B) = \text{Multi}(B)$$

For each next choice of letter, we replace each node from the stacks with its points and compatible children:

```
# stack: list[ChoiceNode]
# letter: char
next_points = points +
sum(n.child[letter].points for n in stacks)
next_stacks = [child for n in stacks for child
in n.child[letter].children]
```

Then:

Theorem: `OrderlyBound` finds all the boards B in a tree with $\text{Multi}(B) \geq S_{\text{high}}$.

Proof: The lemma established an invariant for the recursive calls to `step`. It suffices to check the check the two cases where the function returns early.

If $b < S_{\text{high}}$, then we have:

$$
\begin{aligned}
&b = \text{points} + \sum_{c \in \text{ stack}} \text{Bound}(c) < S_{\text{high}} \\
&\Rightarrow \text{points} + \sum_{c \in \text{ stack}} \text{Force}(c, B) < S_{\text{high}} \forall B \\
&\Rightarrow \text{Multi}(B) < S_{\text{high}} \forall B
\end{aligned}
$$

and therefore there are no high-scoring boards.

If `idx==N`, then the stacks are empty and we have a single board with

$$
\begin{aligned}
b &= \text{points} + \sum_{c \in \text{ stack}} \text{Bound}(c) \\
&= \text{points} \\
&= \text{Multi}(B) \geq S_{\text{high}}
\end{aligned}
$$

So this is a candidate high-scoring board.

Observations:

- OrderlyBound performs a shallow merge.
- OrderlyBound stores at most one pointer to each node in the tree, but in practice many fewer.
- The `points` parameter to `step` is the Multiboggle score on the portion of the board that's been forced.
- OrderlyBound has exponential backtracking behavior. We visit a node in `stack` something like 2^(# of skipped nodes).
- The `branch` operation (tree merging) mitigates this exponential by reducing the number of skips.
- We used a single stack here, but in practice it's much more efficient to keep a separate stack for each cell.

The `branch` and `OrderlyBound` operations work well together. In practice, we build the tree for a board class, then call `branch` some number of times before switching over to `OrderlyBound`. The optimal switchover point is highly variable, but switching when $\text{n.bound} \leq 1.5 S_{\text{high}}$ seems to work well in practice.

TODO: get some stats about how fast OrderlyBound is vs. OrderlyMerge and how much memory they use.

De-duplicated Multiboggle

`OrderlyBound` will report any board with $\text{Multi}(B) \geq S_{\text{high}}$. For each of these, we need to check whether $\text{Score}(B) \geq S_{\text{high}}$ as well. As we saw earlier, these two scores are typically close, but there are some

pathological cases where they diverge. Since `branch` and `OrderlyBound` converge on the Multiboggle score, they'll bog down on the board classes containing these highly duplicative boards, and we'll wind up having to score many millions of Boggle boards to filter them out.

We can improve the situation slightly. Consider the BEE/FEE/BEEF board from earlier:

E	B	E
E	F	E

The max bound counted each of these words four times. If we consider this in the context of a board class, however:

{E,X}	B	{E,X}
{E,X}	F	{E,X}

we can see that "BEE" can be found on the left only when both cells are E, not X, and similarly on the right. But when both of the left cells are E, we know that these paths to BEE, FEE and BEEF will only count once towards the true Boggle score. Both of the left paths to BEE:

```
E-B E   E B E
|       |/
E F E   E F E
```

correspond to the exact same SumNode in the orderly tree. So if we add this word's points to that node once, rather than twice, we don't risk compromising our upper bound.

We do still need to add both the left and the right versions of BEE, FEE and BEEF. If we only added the left versions, then we'd miss the points for this board in the board class:

X	B	E
X	F	E

And we'd no longer have a valid upper bound.

We can call this revised Multiboggle score `DeMulti(B)`. To calculate it, we only score words when they use a unique (unordered) set of cells. So for this board we have:

- `Score(B)` = 3
- `Multi(B)` = 12
- `DeMulti(B)` = 6

Clearly $\text{Score}(B) \leq \text{DeMulti}(B) \leq \text{Multi}(B)$ for all boards B. For boards without repeated letters, $\text{Score}(B) = \text{Multi}(B)$, and so the same holds for $\text{DeMulti}(B)$.

We can filter out duplicate words in `BuildTree`. All the same invariants now hold, only we converge to $\text{DeMulti}(B)$ rather than $\text{Multi}(B)$.

Here are some examples of the effect this deduping has on the root bound for Orderly Trees:

- eeesrvrreeesrsrs: 21953 ⟶ 13253 (true score is 189)
- 4x4 board class containing perslatgsineters: 53037 ⟶ 36881

Since the bound for this board class is so much lower, we expect the Branch and Bound procedure to process it much more quickly. In practice, this is a 5x speedup on this board class. This optimization has the greatest impact on the highest-scoring board classes, which take the greatest time to process.

Final Branch and Bound Algorithm

Here's the final Branch and Bound algorithm for finding Boggle boards B with $\text{Score}(B) \geq S_{\text{high}}$:

1. Enumerate all possible board classes, filtering for symmetry.
2. For each board class C, build an Orderly Tree with deduping.
3. Repeatedly call `branch` until either:
   1. $\text{Bound}(\text{node}) < S_{\text{high}}$ in which case this board class can be eliminated.
   2. $\text{Bound}(\text{node}) \leq 1.5 S_{\text{high}}$ in which case we switch to `OrderlyBound`. This will output a list of boards $B \in C$ such that $\text{DeMulti}(B) \geq S_{\text{high}}$.
4. For each such board B, check whether $\text{Score}(B) \geq S_{\text{high}}$.

This will produce a list of all boards B (up to symmetry) with $\text{Score}(B) \geq S_{\text{high}}$. If two congruent boards fall in the same board class, it will produce both of them.

Results

The Branch and Bound procedure based on Orderly Trees runs significantly faster than the one based on `max_sum_bound`. For 3x3 Boggle with three letter buckets on a single core, the runtime goes from 1h ⟶ 2m, a 30x speedup.

This speedup is greater for larger board classes. Using two letter buckets instead of three reduces the runtime

to just 70s. Compared to the 12 days it would have taken for exhaustive search, this represents a 15,000x speedup.

Complete results for 3x4

Using two letter buckets in the four corners and three buckets for the other eight cells, the Branch and Bound procedure completed in 5h54m on a single core. This represents a 3,000,000x speedup compared to the 2,000 CPU years that exhaustive search would have required.

Using the ENABLE2K wordlist, this search finds 33 distinct boards (up to symmetry) that score 1,500 points or more. Each of these boards can also be found via the hillclimbing procedure, which gives us confidence that it is an effective way to find the global maximum.

Result for one dictionary for 4x4

Two 4x4 runs were completed, one with the ENABLE2K wordlist and one with the NASPA2023 word list. The former was completed before the "deduped Multiboggle" optimization, and its runtime was longer.

- **ENABLE2K**: Found 32 boards with Score≥3500 in 23,000 CPU hours*
- **OSPD5**: Found 46 boards with Score≥3600 in 7,500 CPU hours
- **NASPA2023**: Found 40 boards with Score≥3700 in 9,000 CPU hours.

Compared to exhaustive search, this is roughly a billion times faster. Assuming $0.05/core/hr, this is around $400 of compute.

As with 3x3 and 4x4 Boggle, the top boards can all be found via hillclimbing.

Here are the top five boards for ENABLE2K and NASPA2023:

	ENABLE2K	NASPA2023
1	perslatgsineters: 3625	perslatgsineters: 3923
2	segsrntreiaeslps: 3603	segsrntreiaeslps: 3861
3	gepsnaletireseds: 3593	bestlatepirsseng: 3841
4	aresstapenildres: 3591	dclpeiaerntrsegs: 3835
5	cinslateperidsng: 3591	aresstapenildres: 3826

There is considerable overlap between the highest-scoring boards for each wordlist. NASPA2023 and ENABLE2K share a top board, and top board for OSPD5 is the #2 board for ENABLE2K.

*This run did not include the deduped multiboggle optimization.

Extension to maximizing word count

Instead of seeking the highest-scoring board, we might instead be interested in finding the board with the most words on it. This is a straightforward modification of the problem. We simply change the `SCORE` array to contain all 1s. Then all the tools developed in this paper work exactly as before.

Hillclimbing is also effective at solving this problem. For 3x4 Boggle, the "wordiest" board found via hillclimbing matches the global max found via Branch and Bound. While we haven't run a Branch and Bound search for the wordiest 4x4 board, we'd expect the hillclimbing winner to be the global optimum here as well.

These boards have significant overlap with the highest-scoring boards. The #1 wordiest board for ENABLE2K is serglanepitssero with 1,158 words and 3,569 points. This is also the #8 highest-scoring board. Its middle two columns are identical to those of the highest-scoring board.

S	E	R	G
L	A	N	E
P	I	T	S
S	E	R	O

The "wordiest" known board for ENABLE2K with 1,158 words and 3,569 points.

Variation: Powers of Two Boggle

It's surprising that you get more points for longer words, but only up to eight letters. What if you kept getting more points for longer words? We can set `SCORES=[1, 2, 4, 8, 16, ...]` to play "powers of two" Boggle, where a three letter word is still worth one point, but an eight letter word is worth 32, a sixteen letter word is worth 8,192 points and a seventeen letter word (which must contain a Qu) is worth 16,384 points.

Hillclimbing is *not* effective at finding the highest-scoring board in this version of Boggle. The best board found in 50 hillclimbing runs was `cineqetnsniasesl` (28,542 points). However, we can exhaustively search all boards containing a 17-letter word to find `rpqaselinifcoita` (44,726 points). Over a third of this board's points come from the word "prequalifications."

R	P	Qu	A
S	E	L	I
N	I	F	C
O	I	T	A

The best known board for “powers of two” Boggle, containing the 16,384-point word “prequalifications.”

This failure gives us some insights into why hillclimbing is effective at finding the highest-scoring and wordiest boards. Those scores both produce a smooth fitness landscape, where single character variations on a board produce relatively small changes to its score. This means that the optimal boards are surrounded by other high-scoring boards, and hill climbing is likely to reach the summit if it gets anywhere near it. There’s nowhere for a great board to “hide,” surrounded by mediocre boards.

Contrast this with powers of two Boggle, where a one character change is likely to remove the longest word on a great board.

- highest-scoring neighbor of powers of two: rvqaselinifcoita (20084)
- highest-scoring neighbor of regular winner: perslatgsineders (3488)

Future Work

Branch and Bound with Orderly Trees is extremely effective at exploiting the structure inherent in Boggle, making global searches over the 4x4 board space feasible in a data center. This still takes a lot of compute, however, enough that we haven’t performed full searches for every wordlist, for other languages, or for the most word-dense boards. Further incremental optimizations and reductions in compute costs might make this more attractive.

There is also a 5x5 version of Boggle (”Big Boggle”), and there was a limited run of a 6x6 version (”Super Big Boggle”). These are much harder problems. Even with all the optimizations presented in this paper, an exhaustive search for the best 5x5 Boggle board is expected to take on the order of a billion CPU years. This suggests that a radically different approach is required.

Board Size	CPU time
3x3	70s
3x4	6h
4x4	7500h
4x5	~10,000 years
5x5	~1B years

Here are a few alternative approaches that would be interesting to explore:

- **SMT/ILP solvers** The NP-Hard proof showed that Boggle maximization maps onto well-known SAT problems. These problems are often tackled using ILP Solvers like Z3, OR-Tools or Gurobi. A preliminary investigation didn’t show much promise here, but this is a large field and there may be some better way to frame the problem.
- **GPU acceleration** The process described in this paper relies entirely on the CPU. But the biggest advances in recent years have come from GPUs. It’s not immediately clear how Boggle maximization as described here could be GPU accelerated since the algorithm is tree-heavy and full of data dependencies. Still, there might be another way to frame the problem that results in better acceleration.

L	I	G	D	R
M	A	N	E	S
I	E	T	I	L
D	S	R	A	C
S	E	P	E	S

Best known 5x5 board for ENABLE2K (10,406 points)

Conclusions

Using Branch and Bound, Orderly Trees, and some tailor-made algorithms, we’re able to achieve a factor of a billion speedup over brute-force search. This is enough to prove that the highest-scoring board found via hillclimbing is, in fact, the global optimum. It’s likely that hillclimbing is so effective because the score function produces a relatively smooth fitness landscape. The approach taken in this paper requires solving an NP-Hard problem, and it will not scale well to 5x5 or 6x6 Boggle maximization, which remain well out of reach.